

AW-FNO: An Adaptive approach for Fluid flow Super-resolution via Gated Wavelet-Fourier Learning

Mamta Saini**, Parikshit Mahajan*, Gazania Marine Jyrwa**, Diya Nagchaudhury*, and Sashikumaar Ganesan*,**

Corresponding author: sashi@iisc.ac.in

* Indian Institute of Science, Bengaluru, India.

** ZenteiQ Aitech Innovations Pvt. Ltd.

Keywords: Numerical algorithms, Scientific machine learning, Neural operators, Turbulent flows.

1. Introduction

The study of Partial Differential Equations (PDEs) governs the physical dynamics of a system [1]. Numerical methods solve PDEs with accuracy, but are computationally expensive. Particularly, fluid dynamics simulation at high Reynolds numbers becomes challenging due to the complexity of turbulent interactions. The high computational demand for generating high-fidelity (HF) data through Direct Numerical Simulation (DNS) makes it impractical for real-world applications. As a result, the reconstruction of high-resolution (HR) flow fields from low-resolution (LR) data, called Super Resolution (SR), is a focus of research.

Turbulent flows transfer kinetic energy from coherent large-scale eddies to dissipative small-scale structures. Conventional Convolutional Neural Networks (CNNs) falter due to their spatially local impeding global coherence capture. Neural Operators (NOs) have recently been used for SR paradigms [2] as they learn resolution-invariant mappings between infinite-dimensional function spaces [3]. Fourier Neural Operators (FNOs) [3, 4], leveraging fast Fourier transforms (FFTs) in the Fourier domain, resolve global periodicities but still falter at local discontinuities and sharp gradients.

This work proposes the Augmented wavelet Fourier Neural Operator (AW-FNO), a synergistic architecture fusing the Fourier Neural Operator (FNO) [4] with the Wavelet Neural Operator (WNO) [5]. It combines FNO's spectral efficiency for non-local correlations with WNO's multi-resolution wavelet decomposition. A parallel kernel integration allows hierarchical feature extraction across scales, resulting in superior turbulent flow super-resolution with enhanced generalisation and computational efficiency.

2. Methodology

2.1 The Architecture

We aim to learn an operator $\mathcal{G} : \mathcal{A} \rightarrow \mathcal{U}$ that maps a low-resolution input field $a \in \mathcal{A}$ to a high-resolution output field $u \in \mathcal{U}$.

The core of our approach is the parallel spectral-wavelet layer. The input features are passed through two concurrent paths. The Fourier branch performs a Fast Fourier Transform (FFT), applies a learnable weight matrix to the low-frequency modes, and truncates high-frequency noise. This captures a global structure of the flow. The wavelet branch applies a Discrete Wavelet Transform (DWT) to decompose the signal into different scales. This helps the model learn localised refinements, such as eddies and boundary layer fluctuations.




Figure 1: Architecture of the Augmented Wavelet Fourier Neural Operator

2.2 Mathematical Formulation

The update for a hidden state v_t to v_{t+1} fuses the two branches through a spatially varying gate $\alpha(x)$:

$$v_{t+1}(x) = \sigma \left(Wv_t(x) + \alpha(x) \odot (\mathcal{K}_F v_t)(x) + (1 - \alpha(x)) \odot (\mathcal{K}_W v_t)(x) + b(x) \right), \quad (1)$$

where $\alpha(x) \in [0, 1]$ is the adaptive gate produced by the gated fusion mechanism of Section 2.3, $\odot$ denotes element-wise (channel-wise) multiplication, x is input from the domain, W is a local linear Operator, $\mathcal{K}_F$ is the Fourier kernel operator, defined as $(\mathcal{K}_F(\phi) * (v_t))(x) = \mathcal{F}^{-1}(R_\phi \cdot \mathcal{F}(v_t))(x)$, where $\mathcal{F}$ is Fourier transform, $\mathcal{F}^{-1}$ is inverse Fourier transform, R_ϕ is a neural network-parameterized representation of the kernel k in Fourier space with some parameter $\phi \in \mathcal{R}^P$. $\mathcal{K}_W$ is the Wavelet kernel operator, defined as

$$(\mathcal{K}_W(\phi) * (v_t))(x) = \mathcal{W}^{-1}(R_\phi \cdot \mathcal{W}(v_t))(x) \quad (2)$$

where $\mathcal{W}$ is wavelet transform, $\mathcal{W}^{-1}$ is inverse wavelet transform, and R_ϕ is a neural network-parameterized representation of the kernel k in wavelet space with some parameter $\phi \in \mathcal{R}^P$. b is a bias function and σ is a non-linear activation function.

2.3 Adaptive Spatial Gate

FNO and WNO learn complementary features. The FNO branch operates in the Fourier domain to capture global structures and long-range correlations, whereas the WNO branch uses localised wavelet representations, which excel at representing sharp gradients and multiscale local features. Consequently, no single operator is uniformly optimal across the entire spatial domain: smoothly varying regions and strongly discontinuous regions demand different representational strengths.

To exploit this complementary nature, we use the Adaptive Gated Fusion Mechanism (GFM) [6] that performs spatially adaptive mixing of FNO and WNO features. Let

$$V_{\text{FNO}}, V_{\text{WNO}} \in \mathbb{R}^{B \times H \times W \times C}$$

represent the feature maps produced by FNO and WNO branches, respectively. We first concatenate these features along the channel dimension and pass them through a two-layer convolutional gate with kernel size $k = 5$ and spatial padding, followed by a sigmoid activation, producing a per-channel gating map

$\alpha \in \mathbb{R}^{B \times H \times W \times C}$:

$$\alpha = \sigma \left(\text{Conv}_k \left(\text{GELU} \left(\text{Conv}_k ([V_{\text{FNO}}, V_{\text{WNO}}]) \right) \right) \right).$$

The kernel size is chosen so that the gate’s effective receptive field (a 9-pixel window for $k = 5$ stacked twice) exceeds the characteristic width of the sharp features in the data — approximately 5–10 grid points for Burgers shocks at $\nu = 10^{-3}$. A single 1×1 convolution, having no spatial context, cannot express such a routing pattern and collapses to a near-uniform, indecisive gate; the added depth and spatial receptive field instead let the gate commit decisively to one basis or the other. Whether that commitment is realised as spatial localisation of sharp features or as a per-channel allocation of the two operator sub-spaces is regime-dependent, as analysed in Section 4.3.

The gate α at spatial location induces a piece-wise trust map that encodes the relative importance of FNO versus WNO. The fused representation is then obtained via an element-wise combination,

$$V_{\text{fused}} = \alpha \odot V_{\text{FNO}} + (1 - \alpha) \odot V_{\text{WNO}},$$

where the gating is applied element-wise per channel, allowing different feature channels to route independently. In regions where the underlying solution is smooth and globally coherent, the model learns $\alpha \approx 1$, thereby prioritising the FNO branch and preserving large-scale consistency. Whereas, near the sharp interfaces, discontinuities or highly localised structures, the gate tends to shift to 0, increasing the contribution of WNO and attenuating global spectral effects analogous to the Gibbs phenomenon that arise in purely Fourier-based representations. The GFM provides a router between the two experts, enabling data-dependent choice of expert while yielding a more robust and expressive hybrid operator.

2.4 Entropy-Regularised Gating

Trained with a data-fidelity loss alone, the gate is prone to a “mean-field” stagnation in which it settles at the indecisive equilibrium $\alpha \approx \frac{1}{2}$ everywhere, so that both branches are averaged uniformly and the spatial routing of Section 2.3 never emerges. To break this symmetry, we add a Shannon-entropy penalty on the gate activations to the training objective,

$$L_{\text{total}} = L_{\text{data}} + \lambda_{\text{ent}} \mathbb{E}_x [H(\alpha(x))], \quad H(\alpha) = -\alpha \log \alpha - (1 - \alpha) \log (1 - \alpha), \quad (3)$$

with $\lambda_{\text{ent}} = 0.01$. The penalty drives the per-location gate towards a near-binary (Bernoulli) state, encouraging the model to commit to the locally more appropriate basis — the Fourier branch in smooth, globally coherent regions and the wavelet branch near sharp gradients and discontinuities — rather than averaging the two. Combined with a small random initialisation of the gate convolution (which displaces the gate from the exact $\alpha = \frac{1}{2}$ saddle point), this regulariser produces decisive, spatially varying routing.

3. Experiments

3.1 Dataset

We evaluate AW-FNO on the spatial super-resolution of two-dimensional forced Navier–Stokes turbulence. The data are vorticity snapshots $\omega(x, y)$ of the forced incompressible Navier–Stokes equations on the periodic torus $[0, 2\pi)^2$ at a viscosity corresponding to Reynolds number $\text{Re} \approx 10^3$, providing a turbulent field rich in multi-scale coherent structures. The high-resolution (HR) target is the 128×128 snapshot; the low-resolution (LR) input is obtained by $4\times$ average-pooling to 32×32 followed by bicubic upsampling back to 128×128 , so the network learns the LR→HR residual on a fixed grid. We use 8,000 snapshots for training and 2,000 for testing, with fields normalised to zero mean and unit variance using statistics computed on the training set.

To probe a second, qualitatively different turbulence regime, we also evaluate on the Johns Hopkins Tur-

bulence Database isotropic-turbulence flow (isotropic4096), processed through the identical $128^2/4 \times$ SR pipeline (Section 4.4); unlike the modest-Reynolds forced-NS field, it is strongly intermittent. We additionally use the PDEBench Burgers ($\nu = 10^{-3}$) and 1D shock-tube (Riemann) flows to study the gate on data with genuine discontinuities.

3.2 Baselines and Ablation Variants

We compare AW-FNO against the following models, all trained under an identical optimisation protocol (Section 3.4):

- **Bicubic** — the no-model baseline; the relative error of the bicubic-upsampled LR input itself, quantifying the difficulty of the task.
- **FNO** [4] — a standard Fourier Neural Operator with 4 layers, 32 hidden channels and 20×20 retained Fourier modes.
- **FNO-fat** — a parameter-augmented FNO with 128 hidden channels. This isolates whether additional spectral capacity alone closes the gap to AW-FNO, addressing the fairness of the comparison.
- **WNO** [5] — a Wavelet Neural Operator branch (4 layers, 32 channels, db4 wavelet, decomposition level 3).

To isolate the contribution of the adaptive gate, we evaluate three ablations of AW-FNO that share the same dual-branch backbone (FNO modes 20×20 , WNO db6 at level 3, 32 channels) and differ only in the fusion mechanism:

- **Fixed** $\alpha = 0.5$ — the two branches are averaged with a constant, non-learnable weight, removing all spatial routing.
- **1×1 gate** — the gate is a single 1×1 convolution followed by a sigmoid, i.e. a per-pixel channel mixing with *no* spatial receptive field.
- **Rich gate** ($k = 5$, 2 layers) — the proposed gate of Section 2.3, with a 9×9 effective receptive field, trained with the entropy regulariser ($\lambda_{\text{ent}} = 0.01$) and random gate initialisation.

3.3 Evaluation Metrics

We report (i) the *relative L_2 error*, our primary accuracy metric; (ii) the *relative H_1 error*, which is sensitive to the accuracy of spatial gradients and hence to sharp features; (iii) the *high-frequency spectral L_2 error*, the relative error restricted to wavenumbers above the LR Nyquist cutoff, which directly measures recovery of the small-scale content that bicubic upsampling discards and that Fourier methods tend to corrupt with Gibbs oscillations; and (iv) the *enstrophy error* $|\mathcal{E}_{\text{pred}} - \mathcal{E}_{\text{true}}|/\mathcal{E}_{\text{true}}$ with $\mathcal{E} = \frac{1}{2} \int \omega^2 dx$, a physically meaningful diagnostic of the predicted vorticity field. To characterise the gate we additionally report its mean Shannon entropy $\mathbb{E}_x[H(\alpha)]$ (lower means more decisive routing) and the Pearson correlation $\rho((1 - \alpha), |\nabla \omega|)$ between the WNO-routing weight and the vorticity-gradient magnitude of the target (positive means the gate selects the wavelet branch at sharp features).

3.4 Implementation Details

All models are trained for 500 epochs with the Adam optimiser (learning rate 10^{-3} , weight decay 10^{-4}), a step-decay schedule that halves the learning rate every 100 epochs, gradient clipping at 1.0, and a batch size of 16. Training minimises the relative L_2 loss; for the rich-gate model the entropy penalty of Eq. (3) is added with $\lambda_{\text{ent}} = 0.01$. All runs use a fixed random seed and are executed on a single 48 GB NVIDIA GPU. The checkpoint with the lowest test relative L_2 is retained for evaluation.

Table 1: Super-resolution results on forced Navier–Stokes 128^2 (best-epoch test metrics; lower is better for all columns except where noted). High- f L_2 is the relative spectral error above the LR Nyquist cutoff. Gate entropy H is reported only for gated variants.

Model	Params	Rel L_2	Rel H_1	High- f L_2	Enstrophy	Gate H
Bicubic (no model)	—	0.4134	—	—	—	—
FNO	0.91M	0.0193	0.152	303.8	6.2e-4	—
FNO-fat ($c=128$)	14.6M	0.0198	0.157	325.3	5.6e-4	—
WNO	81.6M	0.0179	0.159	180.6	9.8e-4	—
AW-FNO (fixed $\alpha=0.5$)	89.3M	0.0210	0.194	265.7	1.3e-3	—
AW-FNO (1×1 gate)	89.3M	0.0134	0.116	173.7	1.4e-3	0.667
AW-FNO (rich gate)	89.3M	0.0090	0.081	111.0	2.1e-3	0.017

4. Results and Discussion

4.1 Super-Resolution Accuracy

Table 1 reports the test-set performance of all models on the 128^2 super-resolution task. Every neural operator improves on the bicubic baseline by more than an order of magnitude in relative L_2 error (from 0.413 down to 0.019 for the plain operators and to 0.0094 for AW-FNO with the rich gate), confirming that the small-scale content removed by $4\times$ downsampling is genuinely recoverable and that the task is non-trivial.

Several findings emerge. (i) Among single operators the wavelet branch is strongest on the small scales: the WNO attains the lowest high-frequency spectral error (180.6 versus 303.8 for the FNO), consistent with its multi-resolution basis being better suited to localised vorticity gradients and with the Gibbs-suppression argument of Section 2.3. (ii) Additional Fourier capacity alone does not help: the parameter-augmented FNO-fat (14.6M parameters) is statistically indistinguishable from the standard FNO (0.0198 versus 0.0193), confirming that the gains below come from the dual-branch fusion rather than from raw model size. (iii) The fusion must be *learned*: a non-adaptive average of the two branches ($\alpha = 0.5$) is in fact worse than either branch alone (0.0210), whereas making the mixing coefficient learnable — even with a spatially-blind 1×1 gate — lowers the error to 0.0134, below both single operators. (iv) Accuracy and interpretability are, however, distinct: the 1×1 gate’s entropy stays near the indecisive value ($H \approx 0.67$, cf. $\ln 2 \approx 0.693$) and its routing is spatially flat (Section 4.3), so its gain stems from a learned per-channel re-weighting rather than from feature-aware spatial routing. Endowing the gate with a spatial receptive field (the rich gate) gives the best pointwise and spectral accuracy in the table — relative $L_2 \approx 0.0090$, a $\sim 50\%$ reduction over the best single operator (WNO) and $\sim 33\%$ over the 1×1 gate, together with the lowest relative H_1 (0.081) and high-frequency spectral error (111.0) — driven by a sharply decisive gate ($H \approx 0.017$). The one exception is the enstrophy error, on which the rich gate trades a small amount of global-quantity fidelity for its large pointwise gain. Crucially, as Section 4.3 shows, this decisiveness is realised *per channel* rather than per location, which reshapes how the gate should be interpreted on homogeneous turbulence.

4.2 Energy Spectrum

To assess physical fidelity beyond point-wise error, we compute the radially averaged energy spectrum $E(|k|)$ for each model and compare it with the ground-truth HR field (Figure 2). All methods reproduce the large-scale (low- $|k|$) energy accurately; the models separate in the high- $|k|$ band above the LR Nyquist cutoff, the content that $4\times$ downsampling removes from the input. There, bicubic interpolation injects substantial spurious energy: the fraction of total energy lying above the cutoff is 7.3×10^{-6} , against 1.2×10^{-8} for the ground truth. Every trained operator suppresses this ringing artefact, and the ordering mirrors the high-frequency L_2 column of Table 1 — the spurious high- k fraction falls from 1.8×10^{-6} (FNO) and 1.1×10^{-6} (WNO) to 1.8×10^{-7} for AW-FNO with the rich gate, an order of magnitude




Figure 2: Radially averaged energy spectrum $E(|k|)$ on the NS-forcing 128^2 SR test set. All operators match the low- $|k|$ energy; above the LR Nyquist cutoff the rich-gate AW-FNO tracks the ground-truth tail most closely and injects the least spurious high-wavenumber energy, whereas bicubic upsampling overshoots it the most.

closer to the ground truth than either single operator. This is direct evidence for the Gibbs-suppression mechanism that motivates the wavelet branch (Section 2.3).

4.3 Gate Routing Analysis

A natural hypothesis, borne out on shock-dominated data, is that the spatial gate routes the wavelet branch to sharp features and the Fourier branch to smooth regions. On a viscous Burgers benchmark ($\nu = 10^{-3}$) the same rich gate indeed develops this behaviour, with $\rho((1 - \alpha), |\nabla u|)$ reaching +0.32 along shock fronts. On the forced-turbulence SR task, however, the picture is different and instructive. Extracting the per-block gate maps of the trained rich-gate model and correlating the WNO weight $(1 - \alpha)$ with the vorticity-gradient magnitude $|\nabla \omega|$ of the target yields a correlation indistinguishable from zero ($\rho \approx -0.004$ on the deepest block; individual blocks and samples span $[-0.06, +0.10]$). The channel-averaged gate is spatially almost flat (standard deviation $\approx 2 \times 10^{-3}$), and overlaying its map on the strongest vorticity gradients (Figure 3b) reveals no correspondence: on this data the gate does *not* route spatially.

The decisiveness reported in Table 1 ($H \approx 0.017$) is instead realised along the *channel* axis. Within every fusion block the 32 gate channels split into two saturated groups — about 12 channels with $\alpha \approx 0$ (assigned to WNO) and 20 with $\alpha \approx 1$ (assigned to FNO), with essentially none in between (Figure 3a) — and each channel’s gate is near-constant in space (per-channel spatial standard deviation $\approx 6 \times 10^{-3}$). The rich gate has therefore learned a hard, near-static partition of the feature channels between the two operator sub-spaces, applied uniformly across the domain. The 1×1 gate is spatially flat for the same reason, but, lacking the capacity to commit, leaves all channels near the indecisive value $\alpha \approx 0.5$ ($H \approx 0.67$); the rich gate’s advantage comes precisely from converting this indecision into a sharp channel allocation.

These observations are best read as a *regime characterization*. When the data contains localised, isolated discontinuities (Burgers and shock-tube flows), the optimal use of the two bases is spatial, and the gate routes accordingly. When the data is statistically homogeneous — forced two-dimensional turbulence, in which sharp gradients are vortex filaments distributed throughout the domain rather than confined to




Figure 3: Gate behaviour of the rich-gate AW-FNO on the NS-forcing SR task. (a) Histogram of per-channel mean gate values in the deepest fusion block: the 32 channels split into a near-binary allocation (~ 12 to WNO at $\alpha \approx 0$, ~ 20 to FNO at $\alpha \approx 1$), which accounts for the low gate entropy. (b) The channel-averaged spatial gate map is essentially uniform and shows no alignment with the strongest vorticity gradients (white contours, top-10% $|\nabla\omega|$): on homogeneous turbulence the gate allocates by channel, not by location.

Table 2: SR on JHTDB isotropic turbulence (128^2 , $4\times$). The rich-gate AW-FNO again attains the lowest error and—unlike on forced NS—its gate routes strongly in space.

Model	Rel L_2	Rel H_1	$\rho((1-\alpha), \nabla u)$
FNO	0.0250	—	—
WNO	0.0247	—	—
AW-FNO (fixed $\alpha=0.5$)	0.0251	—	—
AW-FNO (rich gate)	0.0232	0.325	+0.48

a few fronts — no localised sub-region uniformly prefers one basis, and the gate instead exploits the complementary spectral content of the FNO and WNO sub-spaces through a channel-wise allocation applied everywhere. Both are adaptive expert-selection strategies, but along different axes (space versus channel); the accuracy improvement in Table 1 confirms that, in this regime, channel allocation is the effective one.

4.4 Generalisation to a Second Turbulence Regime: JHTDB

To test whether these findings are specific to the forced-NS data, we repeat the SR experiment on the Johns Hopkins Turbulence Database isotropic-turbulence dataset (*isotropic4096*), a high-Reynolds, strongly *intermittent* flow in which energy concentrates on thin vortex filaments embedded in a smooth background. We use the identical $128^2/4\times$ pipeline and the focused four-model set (Table 2).

The rich gate is again the most accurate model. The decisive difference from the forced-NS task is the gate’s spatial behaviour: the routing correlation rises to $\rho \approx +0.48$ (Figure 4), the strongest of any dataset we study and above the Burgers shock value. High-Reynolds isotropic turbulence is genuinely intermittent—its sharp velocity gradients are localised on vortex filaments—so here, unlike the homogeneous forced-NS field, there *are* localised structures for the gate to route to, and it does.

4.5 Cross-Regime Synthesis

Table 3 collects the rich-gate AW-FNO across the flow regimes we study. A single pattern organises the results: *adaptive spatial routing emerges precisely when the flow contains localised sharp structures*,




Figure 4: Rich-gate AW-FNO on JHTDB. The WNO weight $(1 - \alpha)$ correlates with the velocity-gradient magnitude $|\nabla u|$ of the target ($\rho = +0.58$ for this sample), raising the wavelet contribution along high-strain vortex filaments (white contours: top-10% $|\nabla u|$).

and is absent when it does not. The gate routes in space on viscous Burgers shocks ($\rho = +0.17$) and most strongly on intermittent high-Re isotropic turbulence ($\rho = +0.48$); it does not route in space on the statistically homogeneous forced-NS field ($\rho \approx 0$), where it instead realises the FNO/WNO trade-off through a per-channel allocation (Section 4.3). In every regime the rich-gate AW-FNO is the most (or equal-most) accurate operator, but the *mechanism* by which it uses the two bases is regime-dependent. This reframes “adaptive routing” from a fixed behaviour into a data-driven choice between spatial and channel-wise expert selection.

4.6 Ablations

Single-channel (purely spatial) gate. Can the gate be *forced* to route spatially by removing its channel axis? We trained a variant whose gate emits a single map $\alpha(x, y)$ shared across all feature channels, so its only route to decisiveness is spatial. It is uniformly worse: on forced-NS SR it reaches Rel $L_2 = 0.0099$ (vs. 0.0090) and still does not route ($\rho \approx 0$), and on Burgers it degrades to 0.053 (vs. 0.033) while collapsing to a constant per-block basis choice ($\rho \approx 0$). Spatial routing therefore *emerges from* the per-channel structure rather than being suppressed by it; forcing a single spatial mask removes the flexibility that makes the gate effective.

Gate kernel size. Sweeping the gate kernel on Burgers gives Rel L_2 of 0.0364, 0.0333, 0.0341, 0.0345 for $k = 3, 5, 7, 9$, with routing $\rho = +0.14, +0.17, +0.05, +0.10$. $k = 5$ is the sweet spot for both accuracy and routing; enlarging the receptive field past the characteristic feature width brings no further gain.

Table 3: Rich-gate AW-FNO across regimes. The spatial-routing strength $\rho((1 - \alpha), |\nabla \cdot |)$ tracks the spatial intermittency of the *flow*, not the model.

Dataset	Regime	Rel L_2	ρ	Gate H
Forced NS 128^2 SR	homogeneous turbulence	0.0090	≈ 0	0.017
JHTDB iso. SR	high-Re intermittent	0.0232	+0.48	0.011
Burgers $\nu=10^{-3}$	viscous shocks	0.033	+0.17	0.03

Inference cost. Forward latency at 128^2 is 0.60 ms/sample for FNO, 1.83 for WNO, and 5.08 for the rich-gate AW-FNO—an $\sim 8\times$ overhead over a plain FNO but still well within real-time budgets, for a $\sim 50\%$ error reduction.

Limitation. On the PDEBench 1D shock-tube (Riemann) benchmark—predict the developed state of a random two-state initial condition—all operators plateau near Rel $L_2 \approx 0.30$ and the gate does not route ($\rho \approx 0$); this long-horizon map is too hard for the single-step models used here, and we leave it to future work.

5. Conclusions

This work proposes AW-FNO, an adaptive hybrid neural operator that combines the complementary strengths of Fourier and Wavelet representations for fluid flow super-resolution. With the integration of a parallel architecture with an adaptive gated fusion mechanism, we are able to capture the localised multiscale features. The entropy-regularised gating strategy promotes decisive, spatially varying selection between the two operators, thereby enhancing robustness near sharp gradients and discontinuities, and reducing spectral artefacts, such as Gibbs oscillations.

The AW-FNO introduces a hybrid neural operator framework that synergistically integrates the global spectral efficiency of FNOs with the multi-resolution localisation of WNOs for super-resolving turbulent flows. The fusion of a parallel kernel architecture and an adaptive gated fusion mechanism lets the model select between the two bases in a data-dependent way, and this selection is itself regime-dependent: where the flow has localised sharp structures — shock fronts in Burgers and shock-tube flows, or vortex filaments in high-Reynolds intermittent turbulence — the gate routes *spatially*, assigning the wavelet branch to those structures (the routing correlation reaching $\rho \approx 0.48$ on isotropic turbulence); on statistically homogeneous turbulence, by contrast, it commits to a near-binary *per-channel* allocation of the two operator sub-spaces applied across the domain. Both modes require a gate with sufficient representational capacity — a single 1×1 convolution collapses to an indecisive near-uniform mix — and both reduce FNO’s Gibbs oscillations and spectral bias near local structure, yielding high-fidelity reconstruction (a $\sim 50\%$ relative- L_2 improvement over the best single operator on the super-resolution benchmark) with reduced spurious high-wavenumber energy compared with standalone neural operators.

Experiments demonstrate AW-FNO’s robustness on high-Reynolds turbulent benchmarks, achieving zero-shot super-resolution. Future extensions could be dealing with 3D unstructured meshes via graph integrations, further bridging the gap toward real-time CFD surrogates.

References

- [1] Morton E. Gurtin. *An Introduction to Continuum Mechanics*, volume 158 of *Mathematics in Science and Engineering*. Academic Press, 1982.
- [2] Min Wei and Xuesong Zhang. Super-resolution neural operator. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 18247–18256, 2023.
- [3] Nikola Kovachki, Zongyi Li, Burigede Liu, Kamyar Azizzadenesheli, Kaushik Bhattacharya, Andrew

- Stuart, and Anima Anandkumar. Neural operator: Learning maps between function spaces with applications to PDEs. *Journal of Machine Learning Research*, 24(89):1–97, 2023.
- [4] Zongyi Li, Nikola Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Kaushik Bhattacharya, Andrew Stuart, and Anima Anandkumar. Fourier neural operator for parametric partial differential equations. In *International Conference on Learning Representations (ICLR)*, 2021.
- [5] Tapas Tripura and Souvik Chakraborty. Wavelet neural operator for solving parametric partial differential equations in computational mechanics problems. *Computer Methods in Applied Mechanics and Engineering*, 404:115783, 2023.
- [6] Xiaoyi Liu and Hao Tang. DiffFNO: Diffusion fourier neural operator for arbitrary-scale super-resolution. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.